

Executive Summary

Modern Optimizers: Beyond Adam for Deep Learning

Context

The traditional approach of using Adam/AdamW with a single learning rate is becoming outdated. Modern deep learning training leverages:

- **Matrix-oriented optimizers** (Muon, Shampoo, SOAP) that capture parameter correlations
- **Memory-efficient alternatives** (Lion, Adafactor) that scale to billion-parameter models
- **muP principles** for consistent training dynamics across model scales

This Notebook

We implement and compare two modern optimizers against the AdamW baseline:

Optimizer	Type	Key Innovation	Memory	Best For
AdamW	Baseline	Diagonal preconditioning	2x params	General use
Shampoo	Matrix-oriented	Kronecker-factored full-matrix preconditioning	$O(m^2 + n^2)$	Dense layers, unlimited compute
Adafactor	Memory-efficient	Factorized second moments	$\sim 0.5x$ params	Very large models (T5, etc.)

What You'll Do

1. **Implement Shampoo** - Matrix power computation and Kronecker preconditioning
2. **Implement Adafactor** - Factorized second moment estimation
3. **Compare optimizers** - Training dynamics, speed, and final performance on VAE
4. **Hyperparameter exploration** - Find optimal settings for each optimizer

5. Answer analysis questions - Memory, compute, convergence trade-offs

Key Takeaways (Preview)

- Shampoo can converge faster in epochs but is computationally expensive
- Adafactor uses ~4x less memory than Adam for large weight matrices
- Hyperparameter tuning is crucial—defaults rarely optimal
- Choice depends on your constraints: compute budget, memory, model size

Extended VAE Problem: Modern Optimizers

Beyond Adam: Exploring Shampoo and Adafactor

As deep learning training best practices evolve, the traditional approach of using Adam(W) with a single learning rate is being replaced by more sophisticated approaches:

- **Matrix-oriented optimizers** like Muon and Shampoo that capture parameter correlations
- **Memory-efficient alternatives** like Adafactor that scale to massive models
- **muP (maximal update parameterization)** principles for consistent training dynamics

In this extended problem, you will:

1. Implement **Shampoo** - a matrix preconditioning optimizer using Kronecker factors
2. Implement **Adafactor** - a memory-efficient factorized Adam alternative
3. Compare training dynamics of AdamW vs Shampoo vs Adafactor on VAE training
4. Explore hyperparameter choices for each optimizer

Prerequisites: Complete the main VAE implementation (parts a-d) before starting this section.

```
In [ ]: # Install dependencies
!pip install --no-deps --upgrade IPython
!pip install matplotlib numpy torch torchvision tqdm
```

```
In [ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
```

```
from torchvision import datasets, transforms
import matplotlib.pyplot as plt
import numpy as np
from tqdm import tqdm
import time
from typing import Optional, Tuple, List, Dict
import math

# Set device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")
```

Dataset and VAE Model Setup

We'll use MNIST with a simple VAE architecture for comparing optimizers.

```
In [ ]: # Data loading
transform = transforms.Compose([
    transforms.ToTensor(),
])

train_dataset = datasets.MNIST(root='./data', train=True, download=True,
                               transform=transform)
test_dataset = datasets.MNIST(root='./data', train=False, download=True,
                               transform=transform)

batch_size = 128
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True,
                           num_workers=2)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False,
                           num_workers=2)
```

```
In [ ]: class VAE(nn.Module):
    """Simple VAE for MNIST."""
    def __init__(self, input_dim=784, hidden_dim=400, latent_dim=20):
        super().__init__()
        self.input_dim = input_dim
        self.latent_dim = latent_dim

        # Encoder
        self.fc1 = nn.Linear(input_dim, hidden_dim)
        self.fc_mu = nn.Linear(hidden_dim, latent_dim)
        self.fc_logvar = nn.Linear(hidden_dim, latent_dim)

        # Decoder
        self.fc3 = nn.Linear(latent_dim, hidden_dim)
        self.fc4 = nn.Linear(hidden_dim, input_dim)
```

```

def encode(self, x):
    h = F.relu(self.fc1(x))
    return self.fc_mu(h), self.fc_logvar(h)

def reparameterize(self, mu, logvar):
    std = torch.exp(0.5 * logvar)
    eps = torch.randn_like(std)
    return mu + eps * std

def decode(self, z):
    h = F.relu(self.fc3(z))
    return torch.sigmoid(self.fc4(h))

def forward(self, x):
    x = x.view(-1, self.input_dim)
    mu, logvar = self.encode(x)
    z = self.reparameterize(mu, logvar)
    recon = self.decode(z)
    return recon, mu, logvar

def vae_loss(recon_x, x, mu, logvar):
    """VAE loss = Reconstruction + KL divergence."""
    x = x.view(-1, 784)
    BCE = F.binary_cross_entropy(recon_x, x, reduction='sum')
    KLD = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
    return BCE + KLD

```

Part 1: Implementing Shampoo

Background

Shampoo (Gupta et al., 2018) is a matrix preconditioning optimizer that approximates full-matrix AdaGrad using Kronecker-factored statistics.

For a weight matrix $W \in \mathbb{R}^{m \times n}$ with gradient G :

1. Accumulate Kronecker factors:

$$\begin{aligned} L &\leftarrow \beta L + (1-\beta) G G^T \\ R &\leftarrow \beta R + (1-\beta) G^T G \end{aligned}$$

2. Compute preconditioned update:

$$\text{update} = L^{-1/4} \cdot G \cdot R^{-1/4}$$
3. Apply update:

$$W \leftarrow W - \eta \cdot \text{update}$$

The key insight is that $L^{-1/4}$ and $R^{-1/4}$ "whiten" the gradients, leading to better conditioning.

Problem 1.1: Implement Matrix Power via Eigendecomposition

Complete the `matrix_power` function that computes M^p for symmetric positive semi-definite matrices using eigendecomposition.

```
In [ ]: def matrix_power(M: torch.Tensor, p: float, epsilon: float = 1e-6) ->
torch.Tensor:
    """
    Compute M^p for symmetric positive semi-definite matrix M.

    Uses eigendecomposition: M = V @ diag(eigenvalues) @ V^T
    So M^p = V @ diag(eigenvalues^p) @ V^T

    Args:
        M: Symmetric PSD matrix of shape (n, n)
        p: Power to raise the matrix to (e.g., -0.25 for inverse fourth root)
        epsilon: Small constant for numerical stability

    Returns:
        M^p: Matrix of shape (n, n)
    """

#####
# TODO: Implement matrix power using eigendecomposition
#
# 1. Compute eigenvalues and eigenvectors of M using torch.linalg.eigh
#
# 2. Clamp eigenvalues to be >= epsilon for numerical stability
#
# 3. Raise eigenvalues to power p
#
# 4. Reconstruct matrix: V @ diag(eigenvalues^p) @ V^T
#
#####

file:///Users/azmac98/Downloads/q_vae_modern_optimizers.html
```

```
#####
# END OF YOUR CODE
#
#####
return result
```

```
In [ ]: # Test matrix_power
torch.manual_seed(42)
A = torch.randn(4, 4)
M = A @ A.T + 0.1 * torch.eye(4) # Symmetric PSD

M_inv_sqrt = matrix_power(M, -0.5)
identity_approx = M_inv_sqrt @ M @ M_inv_sqrt

print("M^(-0.5) @ M @ M^(-0.5) should be close to identity:")
print(identity_approx)
print(f"Max error from identity: {(identity_approx -
torch.eye(4)).abs().max():.6f}")
assert (identity_approx - torch.eye(4)).abs().max() < 1e-4, "matrix_power
implementation is incorrect"
```

Problem 1.2: Implement Shampoo Optimizer

Complete the Shampoo optimizer class. Key implementation details:

- For 1D parameters (biases), fall back to standard SGD with momentum
- For 2D+ parameters, use Kronecker-factored preconditioning
- Update preconditioners every `update_freq` steps (computing matrix powers is expensive)

```
In [ ]: class Shampoo(torch.optim.Optimizer):
    """
    Shampoo optimizer with Kronecker-factored preconditioning.

    Reference: https://arxiv.org/abs/1802.09568

    Args:
        params: Model parameters
        lr: Learning rate
        momentum: Momentum factor (applied to preconditioned gradients)
        weight_decay: L2 regularization
        epsilon: Numerical stability constant
        update_freq: How often to update preconditioners (every N steps)
    """

    def __init__(self, params, lr=0.1, momentum=0.0, weight_decay=0.0,
                 epsilon=1e-4, update_freq=1):
```

```

defaults = dict(lr=lr, momentum=momentum, weight_decay=weight_decay,
                epsilon=epsilon, update_freq=update_freq)
super().__init__(params, defaults)
self.step_count = 0

@torch.no_grad()
def step(self, closure=None):
    loss = None
    if closure is not None:
        with torch.enable_grad():
            loss = closure()

    self.step_count += 1

    for group in self.param_groups:
        for p in group['params']:
            if p.grad is None:
                continue

            grad = p.grad
            state = self.state[p]

            # Weight decay
            if group['weight_decay'] != 0:
                grad = grad.add(p, alpha=group['weight_decay'])

            # Handle 1D parameters (biases) with simple SGD + momentum
            if p.dim() == 1:
                if 'momentum_buffer' not in state:
                    state['momentum_buffer'] = torch.zeros_like(p)

                buf = state['momentum_buffer']
                buf.mul_(group['momentum']).add_(grad)
                p.add_(buf, alpha=-group['lr'])
                continue

#####
# TODO: Implement Shampoo update for 2D parameters
#
# 1. Reshape gradient to 2D if needed (for conv layers)
# 2. Initialize L and R preconditioners if not present
# 3. Update L and R with current gradient:
#     L = 0.9 * L + 0.1 * (G @ G.T)
#     R = 0.9 * R + 0.1 * (G.T @ G)
# 4. Every update_freq steps, compute:
#     L_inv = matrix_power(L, -0.25, epsilon)
#     R_inv = matrix_power(R, -0.25, epsilon)
# 5. Compute preconditioned gradient:
#     precond_grad = L_inv @ G @ R_inv
# 6. Apply momentum if specified
# 7. Update parameters
#####

```

```

# Reshape to 2D for consistent handling
original_shape = p.shape
if p.dim() > 2:
    grad_2d = grad.view(grad.size(0), -1)
else:
    grad_2d = grad

m, n = grad_2d.shape

# Initialize state
if 'L' not in state:
    state['L'] = group['epsilon'] * torch.eye(m,
device=p.device, dtype=p.dtype)
    state['R'] = group['epsilon'] * torch.eye(n,
device=p.device, dtype=p.dtype)
    state['L_inv'] = torch.eye(m, device=p.device,
dtype=p.dtype)
    state['R_inv'] = torch.eye(n, device=p.device,
dtype=p.dtype)
    if group['momentum'] > 0:
        state['momentum_buffer'] = torch.zeros_like(grad_2d)

#####
#           YOUR CODE HERE
#####
#####

#####
#           END OF YOUR CODE
#####
#####


return loss

```

Solution for Problem 1.2

In []: # SOLUTION

```

def matrix_power_solution(M: torch.Tensor, p: float, epsilon: float = 1e-6) ->
torch.Tensor:
    """Solution for matrix power."""
    eigenvalues, eigenvectors = torch.linalg.eigh(M)
    eigenvalues = torch.clamp(eigenvalues, min=epsilon)
    eigenvalues_p = eigenvalues ** p
    result = eigenvectors @ torch.diag(eigenvalues_p) @ eigenvectors.T
    return result

class ShampooSolution(torch.optim.Optimizer):
    """Complete Shampoo implementation."""

    def __init__(self, params, lr=0.1, momentum=0.0, weight_decay=0.0,

```

```
epsilon=1e-4, update_freq=1, beta=0.9):
    defaults = dict(lr=lr, momentum=momentum, weight_decay=weight_decay,
                    epsilon=epsilon, update_freq=update_freq, beta=beta)
    super().__init__(params, defaults)
    self.step_count = 0

@torch.no_grad()
def step(self, closure=None):
    loss = None
    if closure is not None:
        with torch.enable_grad():
            loss = closure()

    self.step_count += 1

    for group in self.param_groups:
        for p in group['params']:
            if p.grad is None:
                continue

            grad = p.grad
            state = self.state[p]

            # Weight decay
            if group['weight_decay'] != 0:
                grad = grad.add(p, alpha=group['weight_decay'])

            # Handle 1D parameters with simple SGD + momentum
            if p.dim() == 1:
                if 'momentum_buffer' not in state:
                    state['momentum_buffer'] = torch.zeros_like(p)
                buf = state['momentum_buffer']
                buf.mul_(group['momentum']).add_(grad)
                p.add_(buf, alpha=-group['lr'])
                continue

            # Reshape to 2D
            original_shape = p.shape
            if p.dim() > 2:
                grad_2d = grad.view(grad.size(0), -1)
                p_2d = p.view(p.size(0), -1)
            else:
                grad_2d = grad
                p_2d = p

            m, n = grad_2d.shape

            # Initialize state
            if 'L' not in state:
                state['L'] = group['epsilon'] * torch.eye(m,
device=p.device, dtype=p.dtype)
                state['R'] = group['epsilon'] * torch.eye(n,
device=p.device, dtype=p.dtype)
```

```

state['L_inv'] = torch.eye(m, device=p.device,
                           dtype=p.dtype)
state['R_inv'] = torch.eye(n, device=p.device,
                           dtype=p.dtype)
if group['momentum'] > 0:
    state['momentum_buffer'] = torch.zeros_like(grad_2d)

# Update preconditioner statistics
beta = group['beta']
state['L'] = beta * state['L'] + (1 - beta) * (grad_2d @
grad_2d.T)
state['R'] = beta * state['R'] + (1 - beta) * (grad_2d.T @
grad_2d)

# Recompute inverse fourth roots periodically
if self.step_count % group['update_freq'] == 0:
    state['L_inv'] = matrix_power_solution(state['L'], -0.25,
group['epsilon'])
    state['R_inv'] = matrix_power_solution(state['R'], -0.25,
group['epsilon'])

# Compute preconditioned gradient
precond_grad = state['L_inv'] @ grad_2d @ state['R_inv']

# Apply momentum
if group['momentum'] > 0:

state['momentum_buffer'].mul_(group['momentum']).add_(precond_grad)
    update = state['momentum_buffer']
else:
    update = precond_grad

# Update parameters
if p.dim() > 2:
    p.add_(update.view(original_shape), alpha=-group['lr'])
else:
    p.add_(update, alpha=-group['lr'])

return loss

```

Part 2: Implementing Adafactor

Background

Adafactor (Shazeer & Stern, 2018) is a memory-efficient optimizer designed for training large models. It was used to train T5 and other large language models.

Key ideas:

1. **Factorized second moments:** Instead of storing full $m \times n$ second moment matrix, store row and column statistics ($O(m+n)$ vs $O(mn)$)
2. **No momentum** by default (saves memory)
3. **Relative step sizing:** Automatically scale learning rate based on parameter RMS

For a weight matrix $W \in \mathbb{R}^{m \times n}$ with gradient G :

$$\begin{aligned} R_t &= \rho R_{t-1} + (1-\rho) \text{mean}(G^2, \text{dim}=1) \quad \text{(row statistics)} \\ C_t &= \rho C_{t-1} + (1-\rho) \text{mean}(G^2, \text{dim}=0) \quad \text{(column statistics)} \\ \hat{V}_t &= \frac{R_t \otimes C_t}{\text{mean}(R_t)} \quad \text{(reconstructed second moment)} \\ \text{update} &= \frac{G}{\sqrt{\hat{V}_t} + \epsilon} \end{aligned}$$

Problem 2.1: Implement Adafactor

```
In [ ]: class Adafactor(torch.optim.Optimizer):
    """
        Adafactor optimizer with factorized second moment estimation.

        Reference: https://arxiv.org/abs/1804.04235

        Args:
            params: Model parameters
            lr: Learning rate (if None, uses relative step sizing)
            beta2_decay: Decay rate for second moment (-0.8 means rho = 1 - t^(-0.8))
            epsilon1: Regularization constant for second moment
            epsilon2: Regularization constant for RMS
            clip_threshold: Gradient clipping threshold
            weight_decay: L2 regularization
            scale_parameter: Whether to use relative step sizing
    """

    def __init__(self, params, lr=None, beta2_decay=-0.8, epsilon1=1e-30,
                 epsilon2=1e-3, clip_threshold=1.0, weight_decay=0.0,
                 scale_parameter=True):
```

```

        if lr is None and not scale_parameter:
            raise ValueError("lr must be specified if scale_parameter is
False")

        defaults = dict(lr=lr, beta2_decay=beta2_decay, epsilon1=epsilon1,
                         epsilon2=epsilon2, clip_threshold=clip_threshold,
                         weight_decay=weight_decay,
        scale_parameter=scale_parameter)
        super().__init__(params, defaults)

    def _get_rho(self, step, beta2_decay):
        """Compute adaptive decay rate: rho = 1 - step^(beta2_decay)"""
        return min(1.0 - math.pow(step, beta2_decay), 0.999)

    def _rms(self, tensor):
        """Root mean square of tensor."""
        return tensor.norm() / (tensor.numel() ** 0.5)

    @torch.no_grad()
    def step(self, closure=None):
        loss = None
        if closure is not None:
            with torch.enable_grad():
                loss = closure()

        for group in self.param_groups:
            for p in group['params']:
                if p.grad is None:
                    continue

                grad = p.grad
                state = self.state[p]

                # State initialization
                if len(state) == 0:
                    state['step'] = 0

                # For 2D+ parameters, use factorized representation
                if p.dim() >= 2:
                    # Row and column factors
                    state['exp_avg_sq_row'] = torch.zeros(p.shape[:-1],
device=p.device, dtype=p.dtype)
                    state['exp_avg_sq_col'] = torch.zeros(p.shape[:-2] +
p.shape[-1:], device=p.device, dtype=p.dtype)
                else:
                    # For 1D, use full second moment (like Adam)
                    state['exp_avg_sq'] = torch.zeros_like(p)

                state['step'] += 1
                step = state['step']

                # Weight decay
                if group['weight_decay'] != 0:

```

```

grad = grad.add(p, alpha=group['weight_decay'])

rho = self._get_rho(step, group['beta2_decay'])

#####
# TODO: Implement Adafactor update
#
# For 2D+ parameters:
# 1. Compute row mean of grad^2: mean(grad^2, dim=-1)
# 2. Compute col mean of grad^2: mean(grad^2, dim=-2)
# 3. Update exp_avg_sq_row and exp_avg_sq_col with rho
# 4. Reconstruct variance: row[:, None] * col[None, :] /
#                                mean(row)
# 5. Compute update: grad / (sqrt(variance) + epsilon1)
#
# For 1D parameters:
# - Use standard Adam-style second moment
#
# Apply gradient clipping and relative step sizing as needed
#####

#####
#           END OF YOUR CODE
#####

return loss

```

Solution for Problem 2.1

```

In [ ]: # SOLUTION

class AdafactorSolution(torch.optim.Optimizer):
    """Complete Adafactor implementation."""

    def __init__(self, params, lr=None, beta2_decay=-0.8, epsilon1=1e-30,
                 epsilon2=1e-3, clip_threshold=1.0, weight_decay=0.0,
                 scale_parameter=True):

        if lr is None and not scale_parameter:
            raise ValueError("lr must be specified if scale_parameter is
False")

        defaults = dict(lr=lr, beta2_decay=beta2_decay, epsilon1=epsilon1,
                        epsilon2=epsilon2, clip_threshold=clip_threshold,
                        weight_decay=weight_decay,
        scale_parameter=scale_parameter)
        super().__init__(params, defaults)

    def _get_rho(self, step, beta2_decay):
        return min(1.0 - math.pow(step, beta2_decay), 0.999)

```

```

def _rms(self, tensor):
    return tensor.norm() / (tensor.numel() ** 0.5)

def _approx_sq_grad(self, exp_avg_sq_row, exp_avg_sq_col):
    """Reconstruct factorized second moment approximation."""
    r_factor = (exp_avg_sq_row / exp_avg_sq_row.mean(dim=-1,
keepdim=True)).unsqueeze(-1)
    c_factor = exp_avg_sq_col.unsqueeze(-2)
    return r_factor * c_factor

@torch.no_grad()
def step(self, closure=None):
    loss = None
    if closure is not None:
        with torch.enable_grad():
            loss = closure()

    for group in self.param_groups:
        for p in group['params']:
            if p.grad is None:
                continue

            grad = p.grad
            state = self.state[p]
            grad_shape = grad.shape
            factored = len(grad_shape) >= 2

            # State initialization
            if len(state) == 0:
                state['step'] = 0

                if factored:
                    state['exp_avg_sq_row'] = torch.zeros(grad_shape[:-1],
device=p.device, dtype=p.dtype)
                    state['exp_avg_sq_col'] = torch.zeros(grad_shape[:-2]
+ grad_shape[-1:], device=p.device, dtype=p.dtype)
                else:
                    state['exp_avg_sq'] = torch.zeros_like(grad)

            state['RMS'] = 0.0

            state['step'] += 1
            step = state['step']

            # Weight decay (decoupled)
            if group['weight_decay'] != 0:
                p.mul_(1 - group['lr'] * group['weight_decay'] if
group['lr'] else 1 - group['weight_decay'])

            rho = self._get_rho(step, group['beta2_decay'])

            # Update second moment estimate

```

```

if factored:
    # Factorized update
    grad_sqr = grad.pow(2)
    row_mean = grad_sqr.mean(dim=-1)
    col_mean = grad_sqr.mean(dim=-2)

    state['exp_avg_sq_row'].mul_(rho).add_(row_mean, alpha=1 -
    rho)
    state['exp_avg_sq_col'].mul_(rho).add_(col_mean, alpha=1 -
    rho)

    # Reconstruct variance estimate
    variance = self._approx_sq_grad(state['exp_avg_sq_row'],
state['exp_avg_sq_col'])
    variance = variance.clamp(min=group['epsilon1'])
else:
    # Standard (non-factorized) update for 1D params
    state['exp_avg_sq'].mul_(rho).addcmul_(grad, grad, value=1 -
    rho)
    variance =
state['exp_avg_sq'].clamp(min=group['epsilon1'])

    # Compute update
    update = grad / variance.sqrt()

    # Gradient clipping
    update_rms = self._rms(update)
    if update_rms > group['clip_threshold']:
        update.mul_(group['clip_threshold'] / update_rms)

    # Compute learning rate
    if group['lr'] is not None:
        lr = group['lr']
    else:
        # Relative step sizing
        lr = max(group['epsilon2'], self._rms(p))

    if group['scale_parameter'] and group['lr'] is None:
        # Scale by 1/sqrt(step) for warmup
        lr = lr * max(1e-3, 1.0 / math.sqrt(step))

    p.add_(update, alpha=-lr)

return loss

```

Part 3: Optimizer Comparison Experiments

Now we'll compare AdamW, Shampoo, and Adafactor on VAE training.

Problem 3.1: Training Loop and Comparison

```
In [ ]: def train_vae(model, optimizer, train_loader, epochs, device):
    """Train VAE and return training history."""
    model.train()
    history = {'loss': [], 'time_per_epoch': []}

    for epoch in range(epochs):
        epoch_loss = 0
        start_time = time.time()

        for batch_idx, (data, _) in enumerate(train_loader):
            data = data.to(device)
            optimizer.zero_grad()

            recon_batch, mu, logvar = model(data)
            loss = vae_loss(recon_batch, data, mu, logvar)

            loss.backward()
            optimizer.step()

            epoch_loss += loss.item()

        epoch_time = time.time() - start_time
        avg_loss = epoch_loss / len(train_loader.dataset)
        history['loss'].append(avg_loss)
        history['time_per_epoch'].append(epoch_time)

        if (epoch + 1) % 5 == 0:
            print(f"Epoch {epoch+1}/{epochs}, Loss: {avg_loss:.4f}, Time: {epoch_time:.2f}s")

    return history

def evaluate_vae(model, test_loader, device):
    """Evaluate VAE on test set."""
    model.eval()
    test_loss = 0

    with torch.no_grad():
        for data, _ in test_loader:
            data = data.to(device)
            recon, mu, logvar = model(data)
            test_loss += vae_loss(recon, data, mu, logvar).item()
```

```
    return test_loss / len(test_loader.dataset)
```

```
In [ ]: # Run comparison experiment
torch.manual_seed(42)
epochs = 20

results = {}

# AdamW baseline
print("=" * 50)
print("Training with AdamW")
print("=" * 50)
model_adamw = VAE().to(device)
optimizer_adamw = torch.optim.AdamW(model_adamw.parameters(), lr=1e-3,
weight_decay=1e-4)
results['AdamW'] = train_vae(model_adamw, optimizer_adamw, train_loader,
epochs, device)
results['AdamW']['test_loss'] = evaluate_vae(model_adamw, test_loader, device)
print(f"Final test loss: {results['AdamW']['test_loss']:.4f}\n")

# Shampoo
print("=" * 50)
print("Training with Shampoo")
print("=" * 50)
torch.manual_seed(42)
model_shampoo = VAE().to(device)
optimizer_shampoo = ShampooSolution(model_shampoo.parameters(), lr=0.1,
momentum=0.9,
                           weight_decay=1e-4, update_freq=10)
results['Shampoo'] = train_vae(model_shampoo, optimizer_shampoo, train_loader,
epochs, device)
results['Shampoo']['test_loss'] = evaluate_vae(model_shampoo, test_loader,
device)
print(f"Final test loss: {results['Shampoo']['test_loss']:.4f}\n")

# Adafactor
print("=" * 50)
print("Training with Adafactor")
print("=" * 50)
torch.manual_seed(42)
model_adafactor = VAE().to(device)
optimizer_adafactor = AdafactorSolution(model_adafactor.parameters(), lr=1e-2,
weight_decay=1e-4,
scale_parameter=False)
results['Adafactor'] = train_vae(model_adafactor, optimizer_adafactor,
train_loader, epochs, device)
results['Adafactor']['test_loss'] = evaluate_vae(model_adafactor, test_loader,
device)
print(f"Final test loss: {results['Adafactor']['test_loss']:.4f}\n")
```

```
In [ ]: # Plot comparison
fig, axes = plt.subplots(1, 3, figsize=(15, 4))

# Loss curves
ax = axes[0]
for name, data in results.items():
    ax.plot(data['loss'], label=name, linewidth=2)
ax.set_xlabel('Epoch')
ax.set_ylabel('Training Loss')
ax.set_title('Training Loss Comparison')
ax.legend()
ax.grid(True, alpha=0.3)

# Time per epoch
ax = axes[1]
names = list(results.keys())
avg_times = [np.mean(results[n]['time_per_epoch']) for n in names]
ax.bar(names, avg_times, color=['tab:blue', 'tab:orange', 'tab:green'])
ax.set_ylabel('Average Time per Epoch (s)')
ax.set_title('Computational Cost')

# Final test loss
ax = axes[2]
test_losses = [results[n]['test_loss'] for n in names]
ax.bar(names, test_losses, color=['tab:blue', 'tab:orange', 'tab:green'])
ax.set_ylabel('Test Loss')
ax.set_title('Final Test Performance')

plt.tight_layout()
plt.savefig('optimizer_comparison.png', dpi=150)
plt.show()
```

Part 4: Hyperparameter Exploration

Problem 4.1: Shampoo Hyperparameter Study

Explore the effect of key Shampoo hyperparameters:

- `update_freq` : How often to recompute preconditioners (1, 10, 50, 100)
- `lr` : Learning rate (0.01, 0.05, 0.1, 0.5)

```
In [ ]: # Hyperparameter search for Shampoo
shampoo_results = {}

update_freqs = [1, 10, 50]
learning_rates = [0.01, 0.05, 0.1]

for update_freq in update_freqs:
    for lr in learning_rates:
        name = f"freq={update_freq}, lr={lr}"
        print(f"Training Shampoo with {name}")

        torch.manual_seed(42)
        model = VAE().to(device)
        optimizer = ShampooSolution(model.parameters(), lr=lr, momentum=0.9,
                                     update_freq=update_freq)

        history = train_vae(model, optimizer, train_loader, epochs=10,
device=device)
        history['test_loss'] = evaluate_vae(model, test_loader, device)
        shampoo_results[name] = history
        print(f" Final test loss: {history['test_loss']:.4f}")
        print()
```

```
In [ ]: # Visualize Shampoo hyperparameter study
fig, axes = plt.subplots(1, 2, figsize=(12, 4))

# Loss curves
ax = axes[0]
for name, data in shampoo_results.items():
    ax.plot(data['loss'], label=name, linewidth=1.5)
ax.set_xlabel('Epoch')
ax.set_ylabel('Training Loss')
ax.set_title('Shampoo: Learning Rate & Update Frequency Study')
ax.legend(fontsize=8)
ax.grid(True, alpha=0.3)

# Heatmap of final test loss
ax = axes[1]
test_losses = np.zeros((len(update_freqs), len(learning_rates)))
for i, uf in enumerate(update_freqs):
    for j, lr in enumerate(learning_rates):
        name = f"freq={uf}, lr={lr}"
        test_losses[i, j] = shampoo_results[name]['test_loss']

im = ax.imshow(test_losses, cmap='RdYlGn_r')
ax.set_xticks(range(len(learning_rates)))
ax.set_xticklabels([str(lr) for lr in learning_rates])
ax.set_yticks(range(len(update_freqs)))
ax.set_yticklabels([str(uf) for uf in update_freqs])
ax.set_xlabel('Learning Rate')
ax.set_ylabel('Update Frequency')
ax.set_title('Shampoo: Test Loss Heatmap')
```

```

plt.colorbar(im, ax=ax)

# Annotate heatmap
for i in range(len(update_freqs)):
    for j in range(len(learning_rates)):
        ax.text(j, i, f'{test_losses[i,j]:.1f}', ha='center', va='center',
fontsize=9)

plt.tight_layout()
plt.savefig('shampoo_hyperparams.png', dpi=150)
plt.show()

```

Problem 4.2: Adafactor Hyperparameter Study

Explore Adafactor with different settings:

- With and without relative step sizing
- Different learning rates

```

In [ ]: # Hyperparameter search for Adafactor
adafactor_results = {}

configs = [
    {'lr': 1e-3, 'scale_parameter': False, 'name': 'lr=1e-3'},
    {'lr': 5e-3, 'scale_parameter': False, 'name': 'lr=5e-3'},
    {'lr': 1e-2, 'scale_parameter': False, 'name': 'lr=1e-2'},
    {'lr': None, 'scale_parameter': True, 'name': 'relative_step'},
]

for config in configs:
    print(f"Training Adafactor with {config['name']}")

    torch.manual_seed(42)
    model = VAE().to(device)
    optimizer = AdafactorSolution(model.parameters(), lr=config['lr'],
                                  scale_parameter=config['scale_parameter'])

    history = train_vae(model, optimizer, train_loader, epochs=10,
device=device)
    history['test_loss'] = evaluate_vae(model, test_loader, device)
    adafactor_results[config['name']] = history
    print(f" Final test loss: {history['test_loss']:.4f}")
    print()

```

```

In [ ]: # Visualize Adafactor results
fig, axes = plt.subplots(1, 2, figsize=(12, 4))

ax = axes[0]

```

```
for name, data in adafactor_results.items():
    ax.plot(data['loss'], label=name, linewidth=2)
ax.set_xlabel('Epoch')
ax.set_ylabel('Training Loss')
ax.set_title('Adafactor: Learning Rate Study')
ax.legend()
ax.grid(True, alpha=0.3)

ax = axes[1]
names = list(adafactor_results.keys())
test_losses = [adafactor_results[n]['test_loss'] for n in names]
ax.bar(names, test_losses)
ax.set_ylabel('Test Loss')
ax.set_title('Adafactor: Final Test Performance')
ax.tick_params(axis='x', rotation=45)

plt.tight_layout()
plt.savefig('adafactor_hyperparams.png', dpi=150)
plt.show()
```

Part 5: Analysis Questions

Answer the following questions in your writeup:

Question 1: Memory Efficiency

Compare the memory requirements of AdamW, Shampoo, and Adafactor for a weight matrix of shape (m, n) . Which optimizer is most memory-efficient? When might this matter in practice?

Question 2: Computational Trade-offs

Shampoo requires computing matrix inverse fourth roots. Why do we use `update_freq > 1` in practice? How does this affect the trade-off between computational cost and optimization quality?

Question 3: Convergence Properties

Based on your experiments, compare the convergence behavior of the three optimizers. Which converges fastest in terms of epochs? Which is fastest in wall-clock time?

Question 4: Hyperparameter Sensitivity

Based on your hyperparameter exploration, which optimizer seems most sensitive to hyperparameter choices? Which is most robust?

Question 5: Relation to Muon

Explain how Shampoo relates to Muon. Both are "matrix-oriented" optimizers, but they use different techniques. What is the key difference in how they achieve preconditioning?

Summary

In this extended problem, you:

1. **Implemented Shampoo** - a matrix preconditioning optimizer using Kronecker factors
2. **Implemented Adafactor** - a memory-efficient factorized Adam alternative
3. **Compared optimizers** on VAE training, analyzing convergence, speed, and final performance
4. **Explored hyperparameters** for each optimizer

Key takeaways:

- **Shampoo** captures parameter correlations but is computationally expensive
- **Adafactor** is memory-efficient, making it suitable for very large models
- Modern optimizers can outperform Adam, but require careful hyperparameter tuning
- The best optimizer choice depends on model size, compute budget, and task